

# Exploratory Data Analysis using Python - A Case Study

*Analyzing responses from the titanic dataset*

For this lecture we will be working with the [Titanic Data Set from Kaggle](https://www.kaggle.com/c/titanic) (<https://www.kaggle.com/c/titanic>). This is a very famous data set and very often is a student's first step in Data Analytics.

We'll be performing exploratory analysis and answer some hypothesis based on the dataset(at least three).

Some possible hypothesis that we will define for this dataset are:

Let's begin.

- Selecting and downloading a dataset
- Data preparation and cleaning
- Exploratory analysis and visualization
- Asking and answering interesting questions
- Summarizing inferences and drawing conclusions

## Import Libraries

Let's import some libraries to get started!

```
In [19]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

## The Data

Let's start by reading in the titanic\_train.csv file into a pandas dataframe.

Survival value is 1 and 0 indicated not survived

```
In [28]: train = pd.read_csv('titanic_train.csv')
```

```
In [29]: train.head(2)
```

```
Out[29]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...)	female	38.0	1	0	PC 17599	71.2833	C85	S

Let's view some basic information about the data frame.

```
In [30]: train.shape
```

```
Out[30]: (891, 12)
```

The data Frame contains 12 columns and 891 rows

```
In [31]: train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null    int64
1   Survived        891 non-null    int64
2   Pclass          891 non-null    int64
3   Name            891 non-null    object
4   Sex             891 non-null    object
5   Age             714 non-null    float64
6   SibSp           891 non-null    int64
7   Parch           891 non-null    int64
8   Ticket          891 non-null    object
9   Fare            891 non-null    float64
10  Cabin           204 non-null    object
11  Embarked        889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

Most columns have the data type `object`, either because they contain values of different types or contain empty values (`NaN`). It appears that every column contains some empty values since the Non-Null count for every column is lower than the total number of rows (891). We'll need to deal with empty values and manually adjust the data type for each column on a case-by-case basis.

```
In [32]: train['PassengerId'] = pd.to_numeric(train.PassengerId, errors='coerce')
train['Survived'] = pd.to_numeric(train.Survived, errors='coerce')
train['Pclass'] = pd.to_numeric(train.Pclass, errors='coerce')
train['Age'] = pd.to_numeric(train.Age, errors='coerce')
train['SibSp'] = pd.to_numeric(train.SibSp, errors='coerce')
train['Parch'] = pd.to_numeric(train.Parch, errors='coerce')
```

Let's now view some basic statistics about numeric columns.

```
In [33]: train.describe()
```

```
Out[33]:
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
<b>count</b>	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
<b>mean</b>	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
<b>std</b>	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
<b>min</b>	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
<b>25%</b>	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
<b>50%</b>	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
<b>75%</b>	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
<b>max</b>	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

## Exploratory Data Analysis

Let's begin some exploratory data analysis! We'll start by checking out missing data!

### Missing Data

We can use seaborn to create a simple heatmap to see where we are missing data!

```
In [34]: train.isnull()
```


```
Out[34]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Emb
0	False	False	False	False	False	False	False	False	False	False	True	
1	False	False	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	False	True
3	False	False	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	False	True
...	...	...	...	...	...	...	...	...	...	...	...	...
886	False	False	False	False	False	False	False	False	False	False	False	True
887	False	False	False	False	False	False	False	False	False	False	False	False
888	False	False	False	False	False	True	False	False	False	False	False	True
889	False	False	False	False	False	False	False	False	False	False	False	False
890	False	False	False	False	False	False	False	False	False	False	False	True

891 rows × 12 columns

```
In [35]: sns.heatmap(train.isnull(),yticklabels=False,cbar=False,cmap='viridis')
```

```
Out[35]: <AxesSubplot:>
```



Roughly 19 percent of the Age data is missing. The proportion of Age missing is likely small enough for reasonable replacement with some form of imputation. Looking at the Cabin column, it looks like we are just missing too much of that data to do something useful with at a basic level. We'll probably drop this later, or change it to another feature like "Cabin Known: 1 or 0"

In this analysis we wouldn't have to worry much about it

Let's continue on by visualizing some more of the data!

```
In [36]: train.head(3)
```


```
Out[36]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	Na
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C8
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	Na

## Checking the rate of people who survived to those who didn't


```
In [37]: sns.set_style('whitegrid')
sns.countplot(x='Survived', data=train)
```

```
Out[37]: <AxesSubplot:xlabel='Survived', ylabel='count'>
```



## Checking the rate of people who survived based on their gender and those who didn't


```
In [41]: sns.countplot(x='Survived',hue='Sex',data=train,palette='RdBu_r');
```



## Determining the survival rate associated to the class of passengers

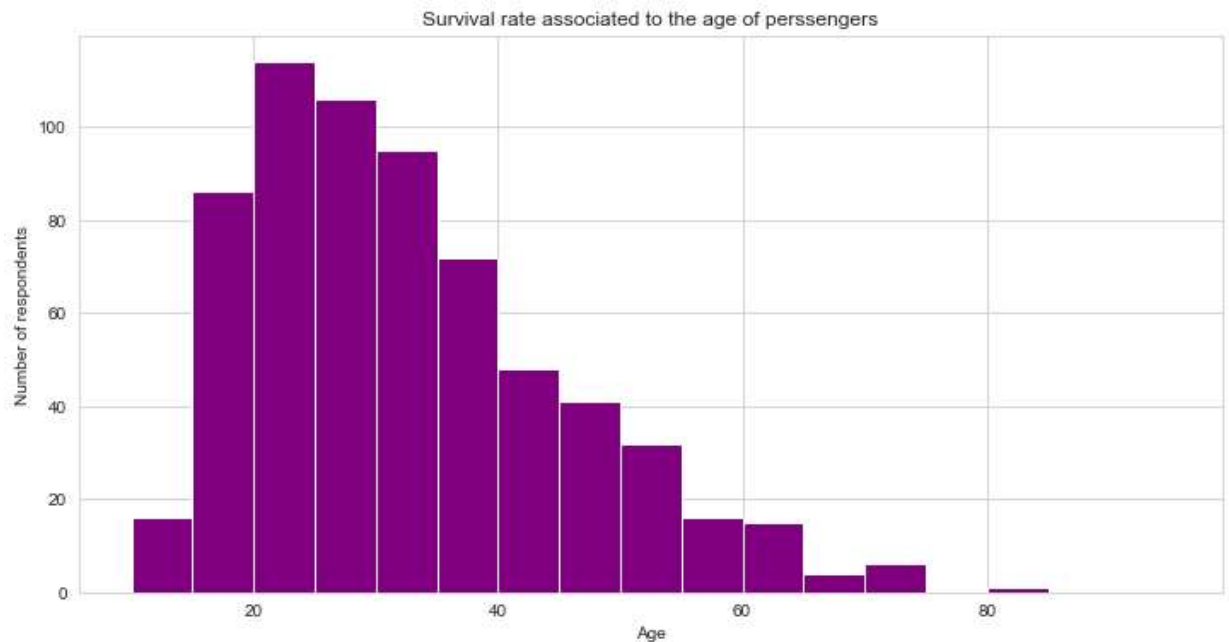
```
In [44]: sns.set_style('whitegrid')  
sns.countplot(x='Survived',hue='Pclass',data=train,palette='rainbow')
```

```
Out[44]: <AxesSubplot:xlabel='Survived', ylabel='count'>
```



```
In [49]: plt.figure(figsize=(12, 6))
plt.title('Survival rate associated to the age of perssengers')
plt.xlabel('Age')
plt.ylabel('Number of respondents')


plt.hist(train.Age, bins=np.arange(10,100,5), color='purple');
```



## Determining the survival rate associated to the age of perssengers

```
In [45]: train['Age'].hist(bins=30,color='darkred',alpha=0.3)
```

```
Out[45]: <AxesSubplot:>
```




## Handling Age Data

We want to fill in missing age data instead of just dropping the missing age data rows. One way to do this is by filling in the mean age of all the passengers (imputation). However we can be smarter about this and check the average age by passenger class. For example:

```
In [26]: plt.figure(figsize=(12, 7))  
sns.boxplot(x='Pclass',y='Age',data=train,palette='winter')
```

```
Out[26]: <matplotlib.axes._subplots.AxesSubplot at 0xe27f780>
```



We can see the wealthier passengers in the higher classes tend to be older, which makes sense. We'll use these average age values to impute based on Pclass for Age.



```
In [27]: def impute_age(cols):
        Age = cols[0]
        Pclass = cols[1]

        if pd.isnull(Age):

            if Pclass == 1:
                return 37

            elif Pclass == 2:
                return 29

            else:
                return 24

        else:
            return Age
```


Now apply that function!

```
In [28]: train['Age'] = train[['Age', 'Pclass']].apply(impute_age,axis=1)
```

Now let's check that heat map again!

```
In [33]: sns.heatmap(train.isnull(),yticklabels=False,cbar=False,cmap='viridis')
```

```
Out[33]: <matplotlib.axes._subplots.AxesSubplot at 0xe4c27b8>
```



## Inferences and Conclusions

We've drawn many inferences from the survey. Here's a summary of a few of them:

- After checking the rate of people who survived to those who didn't, we have been able to see that more life was lost than survived.
- After checking the rate of people who survived based on their gender and those who didn't, we have been able to see that more males died and more females survived.
- After checking the rate of people who survived based on the class of passengers and those who didn't, we have been able to see that more class 3 survived.
- After checking the rate of people who survived based on their age and those who didn't, we have been able to see that people from the age range of 20 - 25 survived more and younger people lost their life same applied to older people from age 58 upwards